

ADITYA NARANJE

+91 9022121737 | Amravati, Maharashtra | aditya.naranje7@gmail.com | linkedin.com/in/anaranje | github.com/adityanaranje

PROFESSIONAL SUMMARY

AI/ML Engineer with 2 years of experience in Python, LangChain, LangGraph, and RAG systems. Experienced in building LLM applications, predictive ML models, and automation pipelines for the insurance domain, with working knowledge of MLOps tools including MLflow and Docker. Strong focus on production AI systems, agentic workflows, evaluation, and measurable automation outcomes.

TECHNICAL SKILLS

AI & NLP	LLM, RAG, Multi-Agent Systems, Prompt Engineering, Embeddings, Transformers, NLP
LLM Providers	OpenAI, Anthropic (Claude), Groq, Google Gemini
Frameworks & Libraries	LangChain, LangGraph, FastAPI, Flask, Streamlit, PyTorch, Scikit-learn, NumPy, Pandas, Matplotlib, Seaborn
Programming	Python, SQL, JavaScript, HTML, CSS
Machine Learning	Supervised & Unsupervised Learning, Classification & Regression, Neural Networks, Model Fine-Tuning, Model Evaluation
MLOps & Tools	MLflow, Docker, Apache Airflow, FAISS, Pinecone, ChromaDB, AWS Bedrock
Platforms & Data	LangSmith, n8n, Hugging Face, Agent Builder, Power BI, SageMaker, AWS Lambda, PostgreSQL/pgvector
Professional	Client Collaboration, Problem Solving, Communication

PROFESSIONAL EXPERIENCE

AI/ML Engineer | Nvest Solutions Private Limited, Thane, Maharashtra

Jun 2024 – Present

- Engineered an LLM-based GitHub PR reviewer using Python, LangGraph, and Flask, automating code validation, severity scoring, and inline fix suggestions; cut manual review effort by ~85% and improved PR turnaround by ~70%.
- Contributed to renewal propensity, churn risk, and lead-scoring ML models using scikit-learn and CatBoost, including Optuna hyperparameter tuning and champion/challenger promotion logic gated on AUC/F1 improvement over rule-based baselines.
- Designed a hybrid policy recommendation engine combining TF-IDF similarity, feature-preference matching, and LightGBM LambdaRank re-ranking, expanding coverage from 1 to 12 lines of business including health, life, motor, and travel.
- Built core components of a production RAG knowledge service using Docling, Voyage AI embeddings, and PostgreSQL/pgvector with hybrid BM25 + vector search, supporting multi-tenant data scoping across insurance clients.
- Collaborated on MLOps observability using MLflow for experiment tracking/model registry, Evidently AI for data-drift monitoring, Prometheus/Grafana dashboards, and Celery automation for daily scoring and weekly retraining.

ServiceNow Internship Trainee | Cognizant, Pune, Maharashtra

Dec 2023 – May 2024

- Designed and implemented an approval workflow and automated notification system in ServiceNow, cutting leave-request processing time by 40% and eliminating manual follow-ups.
- Developed custom business rules, client scripts, and workflow automations to streamline IT service operations, reducing ticket resolution time and manual intervention.
- Built and configured ServiceNow forms and workflows tailored to client requirements, improving process consistency and reducing data-entry errors.

SELECTED PROJECTS

Isolated Multi-Bot RAG System

Python, LangGraph, Pinecone, Docker, AWS EC2, Nginx, Cloudflare

- Developed a production-ready multi-bot RAG platform supporting isolated knowledge bases for 10+ AI bots using Pinecone namespaces.
- Designed a LangGraph decision pipeline to dynamically route queries between vector retrieval and direct LLM responses, improving response relevance by ~35% versus a static RAG pipeline.
- Engineered LLM-driven context relevance scoring (0–1) and answer validation with up to 3 retries, increasing answer accuracy and reducing hallucinations by ~30%.
- Deployed on AWS EC2 with Docker, Nginx, Cloudflare, and Watchtower, reducing manual deployment effort by ~80%.

AI Email Summary Job Agent — Telegram Bot

Gmail API, Telegram Bot API, GitHub Actions, Python

- Developed an AI-powered email automation agent that monitors Gmail for job-related emails and sends structured job alerts to Telegram.
- Implemented OpenAI + LangChain analysis to classify emails, extract company/role/location/experience, and identify required skills such as Python, AWS, SQL, and GenAI.
- Used Gmail API + OAuth2 to fetch unread emails, process content, and mark emails as read to prevent duplicate notifications.
- Automated execution with GitHub Actions every 15 minutes, enabling serverless monitoring without dedicated infrastructure.

RAG Chatbot with Amazon Bedrock AgentCore

Python, Amazon Bedrock AgentCore, LangChain, FAISS

- Developed a fitness Q&A; RAG chatbot using FAISS over 500+ knowledge chunks, improving response relevance by ~40% versus direct LLM responses.
- Integrated Groq Llama-3.3-70B with LangChain tools and deployed the agent on Amazon Bedrock AgentCore runtime.
- Implemented persistent conversation memory using AgentCoreMemoryStore with thread-based cross-session context retention.
- Containerized and deployed through AWS AgentCore CLI to Amazon ECR and Bedrock runtime infrastructure.

EDUCATION

Bachelor of Computer Science — CGPA: 9.14 / 10 2019 – 2023

Prof. Ram Meghe Institute of Technology & Research, Sant Gadge Baba Amravati University

12th Grade, Maharashtra State Board — 85.08% 2018 – 2019

Rajeshwar Union Junior College, Badnera

CERTIFICATIONS & ACHIEVEMENT

Generative AI using LangChain and Hugging Face — Udemy · Introduction to Data Science — IBM · Machine Learning — Internshala · Training Neural Networks in Python — LinkedIn Learning · Feature Engineering — Kaggle

Employee of the Month — July 2025

Recognized for outstanding performance, technical contribution, and consistent delivery of high-quality solutions.

CORE STRENGTHS

- Production Generative AI and RAG engineering
- Agentic workflows, tool calling, and multi-agent orchestration
- ML model development, evaluation, MLOps, and observability
- Cloud deployment, CI/CD, containerization, and scalable APIs
- Business-focused AI automation with measurable outcomes