

Aditya Naranje

AI / ML Engineer | Generative AI | LLM Agents | RAG

aditya.naranje7@gmail.com • Amravati / Maharashtra, India • github.com/adityanaranje • linkedin.com/in/anaranje • adityanaranje.medium.com

PROFESSIONAL SUMMARY

AI/ML Engineer with 2 years of experience designing, deploying, and evaluating production-ready AI, ML, and Generative AI solutions. Skilled in Python, SQL, RAG, LLM agents, agent orchestration, MCP tool integrations, machine-learning pipelines, model evaluation, and AWS deployment. Proven ability to translate business problems into scalable AI systems with measurable gains in automation, reliability, and efficiency.

TECHNICAL SKILLS

AI & GenAI: LLMs, Generative AI, LLM Agents, RAG, Multi-Agent Systems, Prompt Engineering, Embeddings, Vector Search, Tool Calling, Agent Orchestration, Model Evaluation, NLP, Transformers

Orchestration: LangGraph, LangChain, MCP Tool Calling, n8n

Languages & ML: Python, SQL, Pandas, NumPy, Scikit-learn, CatBoost, Optuna, Matplotlib, Seaborn, Regression, Classification

Frameworks & Apps: FastAPI, Flask, Streamlit, LangChain, LangGraph

Cloud & Data: AWS (Bedrock, Lambda, EC2, SageMaker, AgentCore), Docker, Git, GitHub Actions CI/CD, PostgreSQL/pgvector, Pinecone, FAISS

EXPERIENCE

Senior Executive — AI/ML & Generative AI

Nvest Solution Private Limited, Thane, MH • Jul 2024 – Present

- Designed and deployed an LLM-based agent (LangGraph, Flask) for automated GitHub PR review with severity scoring and inline fix suggestions, cutting manual review effort by ~85% and improving PR turnaround by ~70%.
- Built a production RAG knowledge service (Docling parsing, Voyage AI embeddings, PostgreSQL/pgvector with hybrid BM25 + vector search) supporting multi-tenant data scoping across insurance clients.
- Developed a local MCP server for insurance product onboarding with tool-calling/agent orchestration (Claude, Codex, Antigravity) to automate discovery, validation, diagnostics, and SQL generation — cutting setup time ~95% (3–4 days → 2–3 hrs) with read-only DB access and human-in-the-loop review.
- Built ML evaluation & monitoring (renewal-propensity, churn, lead-scoring) with scikit-learn/CatBoost, Optuna tuning and champion/challenger promotion gated on AUC/F1 — expanding recommendation coverage from 1 to 12 insurance lines of business.

SELECTED PROJECTS

Isolated Multi-Bot RAG System (Python, LangGraph, Pinecone, Docker, AWS EC2)

- Production multi-bot RAG platform (10+ bots) with isolated Pinecone namespaces and LangGraph agent routing, improving response relevance ~35%.
- LLM evaluation/validation loop (3-retry, relevance scoring) cut hallucinations ~30%; deployed on AWS EC2 with Docker, Nginx, and Watchtower for automated updates.

AI Email Job Agent — Telegram Bot (Python, Gmail API, Telegram Bot API, GitHub Actions)

- Autonomous agent that monitors Gmail via OAuth2, uses LangChain/OpenAI to classify emails and extract job metadata, and pushes structured alerts to Telegram.
- Shipped fully serverless with scheduled CI/CD via GitHub Actions (15-min cadence), removing dedicated infrastructure.

RAG Chatbot with Amazon Bedrock AgentCore (Python, AWS Bedrock AgentCore, LangChain, FAISS)

- RAG chatbot using FAISS over 500+ knowledge chunks (~40% relevance gain); integrated Groq Llama-3.3-70B with persistent, thread-based memory.
- Containerized and deployed via AWS AgentCore CLI to Amazon ECR / Bedrock runtime.

EDUCATION

Bachelor of Computer Science — CGPA: 9.14 / 10

Prof. Ram Meghe Institute of Technology & Research, SGBAU • 2019 – 2023

CERTIFICATIONS & ACHIEVEMENTS

Generative AI using LangChain & Hugging Face (Udemy) • Data Science (IBM) • Machine Learning (Internshala) • Neural Networks (LinkedIn Learning) • Feature Engineering (Kaggle)